

February 27th, 2019



# Exploiting Nested Task-based Parallelism in the Factorization of Hierarchical Matrices

Rocío Carratalá-Sáez  
Enrique S. Quintana-Ortí

Steffen Börm  
Sven Christophersen

Vicenç Beltran



Christian-Albrechts-Universität zu Kiel



DENSE MATRICES


$\mathcal{H}$  - MATRICES

SPARSE MATRICES

DENSE MATRICES



$\mathcal{H}$  - MATRICES

SPARSE MATRICES




# DENSE MATRICES

# $\mathcal{H}$ - MATRICES



## DENSE MATRICES

## $\mathcal{H}$ - MATRICES



**DENSE  
BLOCKS**



**DENSE  
BLOCKS**

**LOW-RANK  
BLOCKS**

**$O(\log)$  STORAGE [1]**

**$O(\log)$  COMPUTATIONS [1]**

[1] W. Hackbusch. *A sparse matrix arithmetic based on  $\mathcal{H}$ -matrices*. Computing (1999), 62:89-108

# ***H*-LU factorization**

$$\begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} L_{11} & & \\ L_{21} & L_{22} & \\ L_{31} & L_{32} & L_{33} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} & U_{13} \\ & U_{22} & U_{23} \\ & & U_{33} \end{pmatrix}$$

# H-LU factorization

$$\begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} L_{11} & & \\ L_{21} & L_{22} & \\ L_{31} & L_{32} & L_{33} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} & U_{13} \\ & U_{22} & U_{23} \\ & & U_{33} \end{pmatrix}$$

K = 1    K = 2    K = 3



**Algorithm 1** Blocked RL algorithm for the LU factorization.

Require:  $A \in \mathbb{R}^{n \times n}$


```

1: for  $k = 1, 2, \dots, n_t$  do
2:    $A_{kk} = L_{kk} U_{kk}$  LU
3:   for  $j = k + 1, k + 2, \dots, n_t$  do
4:      $U_{kj} := L_{kk}^{-1} A_{kj}$  TRSM
5:   end for
6:   for  $i = k + 1, k + 2, \dots, n_t$  do
7:      $L_{ik} := A_{ik} U_{kk}^{-1}$  TRSM
8:   end for
9:   for  $i = k + 1, k + 2, \dots, n_t$  do
10:    for  $j = k + 1, k + 2, \dots, n_t$  do
11:       $A_{ij} := A_{ij} - L_{ik} \cdot U_{kj}$  GEMM
12:    end for
13:  end for
14: end for
    
```


# H-LU factorization

$$\begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} L_{11} & & \\ L_{21} & L_{22} & \\ L_{31} & L_{32} & L_{33} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} & U_{13} \\ & U_{22} & U_{23} \\ & & U_{33} \end{pmatrix}$$


K = 1







K = 2



K = 3



H-Matrices



**Algorithm 1** Blocked RL algorithm for the LU factorization.

Require:  $A \in \mathbb{R}^{n \times n}$

```

1: for  $k = 1, 2, \dots, n_t$  do
2:    $A_{kk} = L_{kk} U_{kk}$ 
3:   for  $j = k + 1, k + 2, \dots, n_t$  do
4:      $U_{kj} := L_{kk}^{-1} A_{kj}$ 
5:   end for
6:   for  $i = k + 1, k + 2, \dots, n_t$  do
7:      $L_{ik} := A_{ik} U_{kk}^{-1}$ 
8:   end for
9:   for  $i = k + 1, k + 2, \dots, n_t$  do
10:    for  $j = k + 1, k + 2, \dots, n_t$  do
11:       $A_{ij} := A_{ij} - L_{ik} \cdot U_{kj}$ 
12:    end for
13:  end for
14: end for
    
```

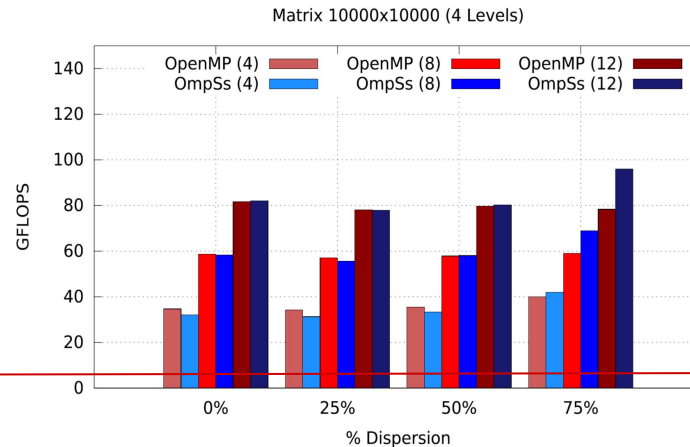

LU

TRSM

GEMM



# $\mathcal{H}$ -LU factorization preliminar results [2]




Sequential  
execution:  
12 GFLOPS

[2] J. I. Aliaga, **R. Carratalá-Sáez**, R. Kriemann, E. S. Quintana-Ortí, *Task-parallel LU factorization of hierarchical matrices using OmpSs*, in: 2017 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW), 2017, pp. 1148–1157. doi:10.1109/IPDPSW.2017.124.



*“There's no need to reinvent the wheel”*





*“There's no need to reinvent the wheel”*



<http://www.h2lib.org>  
Prof. Dr. Steffen Börm  
Christian-Albrechts-Universität zu Kiel

# H2Lib


- Open Source (<https://github.com/H2Lib>)
- C language

# H2Lib

- Open Source (<https://github.com/H2Lib>)
- C language
- $\mathcal{H}$ -Matrices stored as a matrix of (nested) pointers

# H2Lib

- Open Source (<https://github.com/H2Lib>)
- C language
- $\mathcal{H}$ -Matrices stored as a matrix of (nested) pointers
  - 1st limitation: data location (not continuous in memory)



# H2Lib

- Open Source (<https://github.com/H2Lib>)
- C language
- $\mathcal{H}$ -Matrices stored as a matrix of (nested) pointers
  - 1st limitation: data location (not continuous in memory)
- Recursive algorithms


# H2Lib

- Open Source (<https://github.com/H2Lib>)
- C language
- $\mathcal{H}$ -Matrices stored as a matrix of (nested)  $\mathcal{H}$ -matrices
  - 1st limitation: data location (not contiguous)
- Recursive algorithms
  - 2nd limitation: nested dependencies


```
void RECURSIVE_LU( A )
{
    if ( isLeaf( A ) )
        LU( A );
    else
    {
        for ( i = 0 : n_row_sons )
        {
            RECURSIVE_LU( son_ii );
            for ( j = 0 : n_col_sons )
            {
                TRSM( son_ij );
                TRSM( son_ji );
            }
            for ( j = i+1 : n_row_sons )
                for ( k = i+1 : n_col_sons )
                    UPD( son_jk );
        }
    }
}
```



# Data location limitations



# Data location limitations



Forces to partition every leaf block into blocks that present the **SMALLEST** block size **INTERVENING** in the operation.

# Recursion (nested dependencies) limitations

- H2Lib H-LU is **recursive**.
- **Nested** parallelism.

```
void RECURSIVE_LU( A )
{
    if ( isLeaf( A ) )
        LU( A );
    else
    {
        for ( i = 0 : n_row_sons )
        {
            RECURSIVE_LU( son_ii );
            for ( j = 0 : n_col_sons )
            {
                TRSM( son_ij );
                TRSM( son_ji );
            }
            for ( j = i+1 : n_row_sons )
                for ( k = i+1 : n_col_sons )
                    UPD( son_jk );
        }
    }
}
```

# Recursion (nested dependencies) limitations

- H2Lib H-LU is **recursive**.
- **Nested** parallelism.
- **Taskwait** is necessary.
- Parallelism is **limited**.

**TASKWAIT**

```
void RECURSIVE_LU( A )
{
    if ( isLeaf( A ) )
        LU( A );
    else
    {
        for ( i = 0 : n_row_sons )
        {
            RECURSIVE_LU( son_ii );
            for ( j = 0 : n_col_sons )
            {
                TRSM( son_ij );
                TRSM( son_ji );
            }
            for ( j = i+1 : n_row_sons )
                for ( k = i+1 : n_col_sons )
                    UPD( son_jk );
        }
    }
}
```

# The OmpSs \* Programming Model & Runtime

- Task-based programming model
- Data-dependences among tasks expressed with clauses
- Discovers the data-flow parallelism at execution time

(\*) <https://pm.bsc.es/ompss>

# The OmpSs \* Programming Model & Runtime

- Task-based programming model
- Data-dependences among tasks expressed with clauses
- Discovers the data-flow parallelism at execution time
- OmpSs-2 & Nanos6: new features
  - Nested regions dependencies
  - Early release
  - Weak dependencies

(\*) <https://pm.bsc.es/ompss>

## **OmpSs-2 & Nanos6 novelties: nested dependencies**

- Solved OmpSs overlapping PO and regions troubles
- H2Lib limitation: data location
  - Without nested dependencies: partition needed to smallest block size in the matrix.

# OmpSs-2 & Nanos6 novelties: weak dependencies & early release

- H2Lib limitation: recursion
- With weak deps & early release:
  - No taskwait needed.
  - Anticipation of children tasks.

```
#pragma oss task inout(A)
void RECURSIVE_LU( A )
{
    if ( isLeaf( A ) )
        #pragma oss task inout(A)
        LU( A );
    else
    {
        for ( k = 0 : n_row_sons )
        {
            RECURSIVE_LU( son_kk );
            for ( j = 0 : n_col_sons )
            {
                #pragma oss task in(A) inout(son_kj)
                TRSM( son_kj );
                #pragma oss task in(A) inout(son_jk)
                TRSM( son_jk );
            }
            for ( j = k+1 : n_row_sons )
                for ( i = k+1 : n_col_sons )
                    #pragma oss task in(son_ik, son_kj) inout(son_ij)
                    UPD( son_jk );
        }
        #pragma oss taskwait
    }
}
```



# OmpSs-2 & Nanos6 novelties: weak dependencies & early release

- H2Lib limitation: recursion
- With weak deps & early release:
  - No taskwait needed.
  - Anticipation of children tasks.

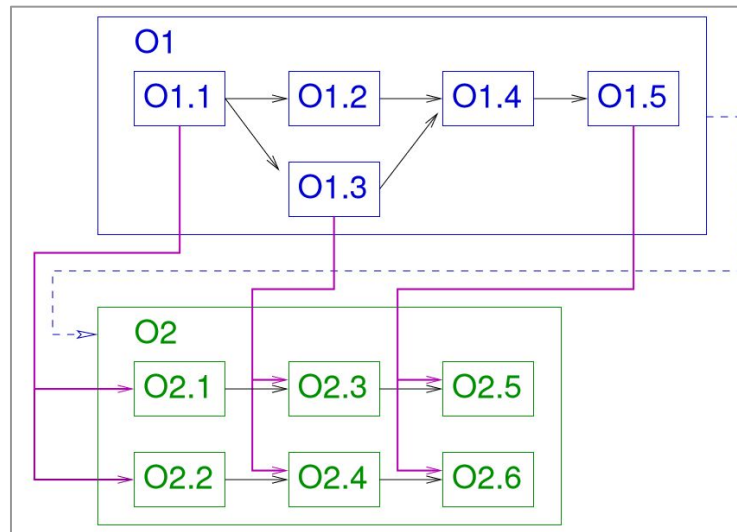
```
#pragma oss task  
weakinout(A)
```

```
#pragma oss task inout(A)  
void RECURSIVE_LU( A )  
{  
    if ( isLeaf( A ) )  
        #pragma oss task inout(A)  
        LU( A );  
    else  
    {  
        for ( k = 0 : n_row_sons )  
        {  
            RECURSIVE_LU( son_kk );  
            for ( j = 0 : n_col_sons )  
            {  
                #pragma oss task in(A) inout(son_kj)  
                TRSM( son_kj );  
                #pragma oss task in(A) inout(son_jk)  
                TRSM( son_jk );  
            }  
            for ( j = k+1 : n_row_sons )  
                for ( i = k+1 : n_col_sons )  
                    #pragma oss task in(son_ik, son_kj) inout(son_ij)  
                    UPD( son_jk );  
        }  
    }  
    #pragma oss taskwait  
}
```

|               |           |           |           |
|---------------|-----------|-----------|-----------|
| $A_{1,1}$     | $A_{1,2}$ | $A_{1,3}$ | $A_{1,4}$ |
| $A_{2,1}$     | $A_{2,2}$ | $A_{2,3}$ | $A_{2,4}$ |
| $A_{3:4,1:2}$ |           | $A_{3,3}$ | $A_{3,4}$ |
|               |           | $A_{4,3}$ | $A_{4,4}$ |

Sequence of operations for the  $\mathcal{H}$ -LU factorization of  $A$ :

|      |   |               |    |   |
|------|---|---------------|----|---|
| O1.1 | : | $A_{1,1}$     | =  | $L_{1,1}U_{1,1}$                              |
| O1.2 | : | $U_{1,2}$     | := | $L_{1,1}^{-1}A_{1,2}$                         |
| O1.3 | : | $L_{2,1}$     | := | $A_{2,1}U_{1,1}^{-1}$                         |
| O1.4 | : | $A_{2,2}$     | := | $A_{2,2} - L_{2,1} \cdot U_{1,2}$             |
| O1.5 | : | $A_{2,2}$     | =  | $L_{2,2}U_{2,2}$                              |
| O2   | : | $U_{1:2,3:4}$ | := | $L_{1:2,1:2}^{-1}A_{1:2,3:4}$                 |
| O3   | : | $L_{3:4,1:2}$ | := | $A_{3:4,1:2}U_{1:2,1:2}^{-1}$                 |
| O4   | : | $A_{3:4,3:4}$ | := | $A_{3:4,3:4} - L_{3:4,1:2} \cdot U_{1:2,3:4}$ |
| O5.1 | : | $A_{3,3}$     | =  | $L_{3,3}U_{3,3}$                              |
| O5.2 | : | $U_{3,4}$     | := | $L_{3,3}^{-1}A_{3,4}$                         |
| O5.3 | : | $L_{4,3}$     | := | $A_{4,3}U_{3,3}^{-1}$                         |
| O5.4 | : | $A_{4,4}$     | := | $A_{4,4} - L_{4,3} \cdot U_{3,4}$             |
| O5.5 | : | $A_{4,4}$     | =  | $L_{4,4}U_{4,4}$                              |



|      |   |           |    |                                    |
|------|---|-----------|----|------------------------------------|
| O2.1 | : | $U_{1,3}$ | := | $L_{1,1}^{-1}A_{1,3}$ ,            |
| O2.2 | : | $U_{1,4}$ | := | $L_{1,1}^{-1}A_{1,4}$ ,            |
| O2.3 | : | $A_{2,3}$ | := | $A_{2,3} - L_{21} \cdot U_{1,3}$ , |
| O2.4 | : | $A_{2,4}$ | := | $A_{2,4} - L_{21} \cdot U_{1,4}$ , |
| O2.5 | : | $U_{2,3}$ | := | $L_{2,2}^{-1}A_{2,3}$ , and        |
| O2.6 | : | $U_{2,4}$ | := | $L_{2,2}^{-1}A_{2,4}$ .            |

# Tests & results

- Integral equations, BEM
- Underlying kernel function: Laplace equation

$$g : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}, \quad g(x, y) = \begin{cases} -\log |x - y| & : d = 1, \\ -\frac{1}{2\pi} \log \|x - y\|_2 & : d = 2, \\ \frac{1}{4\pi} \|x - y\|_2^{-1} & : d = 3. \end{cases}$$


- Compression with SVD.
- Admissibility condition  $\max\{\text{diam}(\mathcal{B}_t), \text{diam}(\mathcal{B}_s)\} \leq \eta \text{dist}(\mathcal{B}_t, \mathcal{B}_s)$  where  $\eta \in \mathbb{R}_{>0}$

## SETUP:

- IEEE 754 double precision arithmetic.
- Single node of the MareNostrum 4 system at BSC, with 2 Intel Xeon Platinum 8160 sockets, 24 cores per socket, 96 Gbytes of DDR4 RAM.
- GCC 4.8.5, Intel MKL 2017.4 (AVX2 instructions enabled), OmpSs-2 (mcxx 2.1.0).


# Tests & results

2D BEM (30 K)



# Tests & results

3D BEM (42K)



# Conclusions

From **previous** [2] work...

- **Task-based parallelism** seems to offer good performance for H-structures.

From **using OmpSs 2** [3] on H2Lib H-LU...

- Good parallel performance with OmpSs 2 **task-based parallelism**.
- **New features (weak dependencies & early release)** avoid detected limitations.

[2] R. Carratalá-Sáez et al. *Task-parallel LU factorization of hierarchical matrices using OmpSs*. Proceedings of the 19th IEEE Workshop on Parallel and Distributed Scientific and Engineering Computing, PDSEC 2017, 2017.

[3] R. Carratalá-Sáez et al. *Exploiting Nested Task-Parallelism in the H-LU Factorization*. Journal of Computational Science (Elsevier), February 2018. Accepted, pending of publication.

## Future work

- Related to H2Lib:
  - Maybe go deeper in parallelism.
  - Study how to improve parallel efficiency in 2D.
  - More test cases (different applications).
  - Other operations.

## Future work

- Related to H2Lib:
  - Maybe go deeper in parallelism.
  - Study how to improve parallel efficiency in 2D.
  - More test cases (different applications).
  - Other operations.
- Implement distributed memory H-LU.
- Chameleon extension (low-rank support).

*Inria*



# Exploiting Nested Task-based Parallelism in the Factorization of Hierarchical Matrices



## Thanks for your attention!



You are more than welcome to contact us:

**Rocío Carratalá-Sáez**

Enrique S. Quintana-Ortí

Steffen Börm

Sven Christophersen

Vicenç Beltrán

**rcarrata@uji.es**

quintana@uji.es

boerm@math.uni-kiel.de

christophersen@math.uni-kiel.de

vbeltran@bsc.es



Christian-Albrechts-Universität zu Kiel

