

# TASK-BASED SPARSE DIRECT SOLVER FOR SYMMETRIC INDEFINITE SYSTEMS

---

Iain S. Duff and **Florent Lopez**

SIAM CSE19, Spokane, WA, 2019

Rutherford Appleton Laboratory

# Objectives

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

# Objectives

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

- Fast

# Objectives

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

- Fast
- Numerically Stable

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

- Fast
- Numerically Stable

In the context of the [NLAfet project](#), we have developed the package [SyLVER](#), that contains the two direct solvers:

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

- Fast
- Numerically Stable

In the context of the NLAfet project, we have developed the package SyLVER, that contains the two direct solvers:

- SpLDLT: for symmetric systems:
  - positive-definite ( $LL^T$ )
  - indefinite ( $LDL^T$ )

Objective: solve  $Ax = b$

Where  $A$  is a large, sparse and symmetrically-structured matrix

The algorithm needs to be:

- Fast
- Numerically Stable

In the context of the NLAfet project, we have developed the package SyLVER, that contains the two direct solvers:

- SpLDLT: for symmetric systems:
  - positive-definite ( $LL^T$ )
  - indefinite ( $LDL^T$ )
- SpLU: for unsymmetric ( $LU$ ) systems

# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:



# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:

Top500<sup>1</sup>: #1 (Power9&V100), #2 (Power9&V100), #5 (Haswell&P100)

---

<sup>1</sup><https://www.top500.org/>

# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:

Top500<sup>1</sup>: #1 (Power9&V100), #2 (Power9&V100), #5 (Haswell&P100)

GPU architectures:

- ▲ High performance peak (GFlop/sec, GFlop/watt)  
e.g. NVIDIA V100 GPU: 7.8 TFlop/s theoretical peak (FP64), Cholesky peak 5.7 TFlop/s (FP64)

---

<sup>1</sup><https://www.top500.org/>

# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:

Top500<sup>1</sup>: #1 (Power9&V100), #2 (Power9&V100), #5 (Haswell&P100)

GPU architectures:

- ▲ High performance peak (GFlop/sec, GFlop/watt)  
e.g. NVIDIA V100 GPU: 7.8 TFlop/s theoretical peak (FP64), Cholesky peak 5.7 TFlop/s (FP64)
- ▲ Support for reduced precision arithmetic such as FP16  
e.g. Cholesky peak 46.0 TFlop/s (FP16 using Tensor Cores)

---

<sup>1</sup><https://www.top500.org/>

# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:

Top500<sup>1</sup>: #1 (Power9&V100), #2 (Power9&V100), #5 (Haswell&P100)

GPU architectures:

- ▲ High performance peak (GFlop/sec, GFlop/watt)  
e.g. NVIDIA V100 GPU: 7.8 TFlop/s theoretical peak (FP64), Cholesky peak 5.7 TFlop/s (FP64)
- ▲ Support for reduced precision arithmetic such as FP16  
e.g. Cholesky peak 46.0 TFlop/s (FP16 using Tensor Cores)
- ▼ Limited memory available on the device (16GB on V100)

---

<sup>1</sup><https://www.top500.org/>

# Heterogenous CPU-GPU systems

Heterogeneous CPU-GPU architectures have been increasingly popular in the HPC community over the past few years:

Top500<sup>1</sup>: #1 (Power9&V100), #2 (Power9&V100), #5 (Haswell&P100)

GPU architectures:

- ▲ High performance peak (GFlop/sec, GFlop/watt)  
e.g. NVIDIA V100 GPU: 7.8 TFlop/s theoretical peak (FP64), Cholesky peak 5.7 TFlop/s (FP64)
- ▲ Support for reduced precision arithmetic such as FP16  
e.g. Cholesky peak 46.0 TFlop/s (FP16 using Tensor Cores)
- ▼ Limited memory available on the device (16GB on V100)
- ▼ CPU ↔ GPU memory transfer potentially slow  
e.g. PCIe 10× slower than NVLink but still common


---

<sup>1</sup><https://www.top500.org/>

# Sparse direct methods: parallelism

The **multifrontal factorization** is achieved with a topological traversal of the **elimination tree**:

1. **Assemble** the contributions from the descendant
2. **Factor** the column associated with the current node
3. **Form** the contribution blocks




# Sparse direct methods: parallelism

The **multifrontal factorization** is achieved with a topological traversal of the **elimination tree**:

1. **Assemble** the contributions from the descendant
2. **Factor** the column associated with the current node
3. **Form** the contribution blocks

For the **factorization** and the **solve** phases, the sources of **parallelism** in the *elimination tree* include:

- **Tree parallelism**: nodes in independent branches can be processed concurrently
- **Node parallelism**: when a node is large enough, it may be processed in parallel




# MULTIFRONTAL FACTORIZATION FOR MULTICORE CPUS


---




# Dense Cholesky factorization on multicore CPUs



# Dense Cholesky factorization on multicore CPUs



# Dense Cholesky factorization on multicore CPUs





**Factor** diagonal block  $A_{kk} = L_{kk} L_{kk}^T$

**Apply** pivots  $L_{ik} = A_{ik} L_{kk}^{-T}$

**Update** trailing sub-matrix  $A_{ij-} = L_{ik} L_{jk}^T$

# Dense Cholesky factorization on multicore CPUs




**Factor** diagonal block  $A_{kk} = L_{kk} L_{kk}^T$

**Apply** pivots  $L_{ik} = A_{ik} L_{kk}^{-T}$

**Update** trailing sub-matrix  $A_{ij-} = L_{ik} L_{jk}^T$

# Dense Cholesky factorization on multicore CPUs




**Indefinite systems:** small diagonal elements and/or large off-diagonal elements might cause a loss of accuracy

**Factor** diagonal block  $A_{kk} = L_{kk} L_{kk}^T$

**Apply** pivots  $L_{ik} = A_{ik} L_{kk}^{-T}$

**Update** trailing sub-matrix  $A_{ij-} = L_{ik} L_{jk}^T$

# Dense Cholesky factorization on multicore CPUs



**Indefinite systems:** small diagonal elements and/or large off-diagonal elements might cause a loss of accuracy

⇒ avoid it by using **pivoting**


**Factor** diagonal block  $A_{kk} = L_{kk} L_{kk}^T$

**Apply** pivots  $L_{ik} = A_{ik} L_{kk}^{-T}$

**Update** trailing sub-matrix  $A_{ij-} = L_{ik} L_{jk}^T$


# Pivoting strategies

Two main pivoting strategies for the  $LDL^T$  factorization:



# Pivoting strategies

Two main pivoting strategies for the  $LDL^T$  factorization:




- **Threshold Partial Pivoting (TPP)**
  - ▲ Ensures that  $|l_{ij}| < u^{-1}$  for some threshold  $u$
  - ▼ Requires global communication on each columns
    - Used in HSL\_MA97

Numerically robust but hard to parallelize



# Pivoting strategies

Two main pivoting strategies for the  $LDL^T$  factorization:



- **Threshold Partial Pivoting (TPP)**
  - ▲ Ensures that  $|l_{ij}| < u^{-1}$  for some threshold  $u$
  - ▼ Requires global communication on each columns
    - Used in HSL\_MA97

Numerically robust but hard to parallelize

- **Supernode Bunch-Kaufmann (SBK)**
  - ▲ Performs pivoting in only within  $F_{11}$  using Bunch-Kaufmann algorithm
  - ▼ Requires the use of a pivots perturbation strategy
    - Scaling and ordering techniques may be used to alleviate this problem
    - Used in PARDISO

Plenty of parallelism but potentially unstable

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization


1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
2. **Failed-in-place approach**: keep the failed columns in place:

# A Posteriori Threshold Pivoting




- Factor diagonal block using complete pivoting

```
do k = 1, nblk  
  ! Factor diagonal block  
  call Factor(A(k,k):RW, nelim_k)
```

```
end do
```

# A Posteriori Threshold Pivoting




- Apply pivots in parallel using speculative execution

```
do k = 1, nblk
  ! Factor diagonal block
  call Factor(A(k,k):RW, nelim_k)
  ! Compute factors sub-diagonal on blocks
  do j = k+1, mblk
    call ApplyN(A(k,j):RW, A(k,k):R, nelim_k)
  end do
  ! Compute factors left-diagonal on blocks
  do j = 1, k-1
    call ApplyT(A(k,j):RW, A(k,k):R, nelim_k)
  end do
```

```
end do
```

## A Posteriori Threshold Pivoting




- Determine the number of **successfully eliminated** pivots w.r.t threshold  $\mu^{-1}$

```
do k = 1, nblk
! Factor diagonal block
call Factor(A(k,k):RW, nelim_k)
! Compute factors sub-diagonal on blocks
do j = k+1, mblk
    call ApplyN(A(k,j):RW, A(k,k):R, nelim_k)
end do
! Compute factors left-diagonal on blocks
do j = 1, k-1
    call ApplyT(A(k,j):RW, A(k,k):R, nelim_k)
end do
! Compute  $nelim_k$ 
call Adjust(A(:,k):R, nelim_k)
```

end do

# A Posteriori Threshold Pivoting




- Update blocks in the trailing sub-matrix

```
do k = 1, nblk
  ! Factor diagonal block
  call Factor(A(k,k):RW, nelim_k)
  ! Compute factors sub-diagonal on blocks
  do j = k+1, mblk
    call ApplyN(A(k,j):RW, A(k,k):R, nelim_k)
  end do
  ! Compute factors left-diagonal on blocks
  do j = 1, k-1
    call ApplyT(A(k,j):RW, A(k,k):R, nelim_k)
  end do
  ! Compute  $nelim_k$ 
  call Adjust(A(:,k):R, nelim_k)
  ! Update trailing sub-matrix
  do j = k, nblk
    do i = j, mblk
      call UpdateNN(A(i,k):R, A(j,k):R, A(i,j):RW,
                    nelim_k)
    end do
  end do
end do
```



# A Posteriori Threshold Pivoting



- Update **uneliminated** and **failed entries** in the left-diagonal blocks

```
do k = 1, nblk
  ! Factor diagonal block
  call Factor(A(k,k):RW, nelim_k)
  ! Compute factors sub-diagonal on blocks
  do j = k+1, mblk
    call ApplyN(A(k,j):RW, A(k,k):R, nelim_k)
  end do
  ! Compute factors left-diagonal on blocks
  do j = 1, k-1
    call ApplyT(A(k,j):RW, A(k,k):R, nelim_k)
  end do
  ! Compute  $nelim_k$ 
  call Adjust(A(:,k):R, nelim_k)
  ! Update trailing sub-matrix
  do j = k, nblk
    do i = j, mblk
      call UpdateNN(A(i,k):R, A(j,k):R, A(i,j):RW,
                    nelim_k)
    end do
  end do
  ! Update uneliminated entries on left-diagonal
  ! blocks
  do j = 1, k-1
    do i = j, k-1
      call UpdateTT(A(k,i):R, A(k,j):R, A(i,j):RW,
                    nelim_k)
    end do
    do i = k, mblk
      call UpdateTT(A(i,k):R, A(k,j):R, A(i,j):RW,
                    nelim_k)
    end do
  end do
end do
```

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
2. **Failed-in-place approach**: keep the failed columns in place:

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
  - ▲ **Increased parallelism**
2. **Failed-in-place approach**: keep the failed columns in place:

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
  - ▲ **Increased parallelism**
  - ▼ Backups of entries  $\implies$  **increased memory footprint**
2. **Failed-in-place approach**: keep the failed columns in place:

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
  - ▲ Increased parallelism
  - ▼ Backups of entries  $\implies$  increased memory footprint
2. **Failed-in-place approach**: keep the failed columns in place:
  - ▲ Dramatically reduce data movement

# A Posteriori Threshold Pivoting

A Posteriori Threshold Pivoting (APTP) strategy: improve the **parallelism** of TPP while preserving **numerical stability** of the factorization

1. **Speculative execution**: execute some tasks speculatively, assuming no numerical issues occurred, then check for instability and backtrack if necessary:
  - ▲ Increased parallelism
  - ▼ Backups of entries  $\implies$  increased memory footprint
2. **Failed-in-place approach**: keep the failed columns in place:
  - ▲ Dramatically reduce data movement
  - ▼ Failed entries must be kept up-to-date  $\implies$  Reduced granularity

# Parallel multifrontal factorization for multicore CPUs

SpLDLT implements a multifrontal factorization using a Sequential Task Flow (STF) model

```
do n = 1, nnodes ! Topologically ordered
  ! Allocate memory
  call activate(front(n));
  ! Assemble fully-summed columns
  call assemble(front(n), children(n));
  ! Compute factors
  call factor(front(n));
  ! Assemble contribution block
  call assemble_contrib(front(n), children(n));
end do

! Wait for completion of submitted tasks
call task_wait_for_all();
```


- **assemble**
  1. Allocate memory (fully-summed and contribution blocks)
  2. Assemble contributions from children nodes into fully-summed columns
- **factor** Factorize the fully-summed columns and form the contribution blocks
- **assemble\_contrib** Assemble contributions from children nodes into contribution blocks

# MULTIFRONTAL FACTORIZATION FOR GPU DEVICES


---




# Multifrontal factorization for GPU devices



# Multifrontal factorization for GPU devices




# Multifrontal factorization for GPU devices



```
do lvl = 1, nlevels
  ! Allocate memory
  call activate_lvl(fronts(lvl))
  ! Assemble fully-summed columns
  call assemble_lvl(fronts(lvl), child(lvl))
  ! Compute factors
  call factor_lvl(fronts(lvl))
  ! Assemble contribution block
  call assemble_cb_lvl(fronts(lvl), child(lvl))
end do
```

# Multifrontal factorization for GPU devices




```
do lvl = 1, nlevels
  ! Allocate memory
  call activate_lvl(fronts(lvl))
  ! Assemble fully-summed columns
  call assemble_lvl(fronts(lvl), child(lvl))
  ! Compute factors
  call factor_lvl(fronts(lvl))
  ! Assemble contribution block
  call assemble_cb_lvl(fronts(lvl), child(lvl))
end do
```

```
chol_gpu: do jj = 1, n(lvl), nb
  do kk = jj, jj+nb, ib
    ! Factor inner panel
    call potrf_gpu_batched(lvlnodes(kk))
    ! Update inner panel
    call syrk_gpu_batched(lvlnodes(kk))
    ! Update trailing submatrix
  end do
  call syrk_gpu_batched(lvlnodes(jj), &
    lvlnodes(jj+nb))
end do chol_gpu
```

# HETEROGENEOUS CPU-GPU MULTIFRONTAL FACTORIZATION


---

# Heterogeneous CPU-GPU multifrontal factorization




# Heterogeneous CPU-GPU multifrontal factorization

- Partition the assembly tree between NUMA regions and GPU devices




# Heterogeneous CPU-GPU multifrontal factorization



- Partition the assembly tree between NUMA regions and GPU devices
- $\text{balance} = \frac{\max_i(w_i)}{\frac{1}{n_{\text{res}}} \sum_j (w_j)}$




# Heterogeneous CPU-GPU multifrontal factorization




- Partition the assembly tree between NUMA regions and GPU devices
- $\text{balance} = \frac{\max_i(w_i/\alpha_i)}{\frac{1}{n_{\text{res}}} \sum_j (w_j/\alpha_j)}$ 
  - $\alpha_i = 1.0$  for NUMA regions
  - $\alpha_i = \text{Perf}_{\text{GPU}}/\text{Perf}_{\text{CPU}}$  for GPU devices

# Heterogeneous CPU-GPU multifrontal factorization




- Partition the assembly tree between NUMA regions and GPU devices
- $\text{balance} = \frac{\max_i(w_i/\alpha_i)}{\frac{1}{n_{\text{res}}} \sum_j (w_j/\alpha_j)}$ 
  - $\alpha_i = 1.0$  for NUMA regions
  - $\alpha_i = \text{Perf}_{\text{GPU}}/\text{Perf}_{\text{CPU}}$  for GPU devices
- Within subtree:
  - NUMA partition: task-based factorization
  - GPU partition: batched factorization

# Heterogeneous CPU-GPU multifrontal factorization



- Partition the assembly tree between NUMA regions and GPU devices
- $$\text{balance} = \frac{\max_i(w_i/\alpha_i)}{\frac{1}{n_{\text{res}}} \sum_j (w_j/\alpha_j)}$$
  - $\alpha_i = 1.0$  for NUMA regions
  - $\alpha_i = \text{Perf}_{\text{GPU}}/\text{Perf}_{\text{CPU}}$  for GPU devices
- Within subtree:
  - NUMA partition: task-based factorization
  - GPU partition: batched factorization
- Root partition: use task-based factorization where tasks run either on a CPU core or in a GPU stream

# Heterogeneous CPU-GPU multifrontal factorization



- Partition the assembly tree between NUMA regions and GPU devices
- $\text{balance} = \frac{\max_i(w_i/\alpha_i)}{\frac{1}{n_{\text{res}}} \sum_j (w_j/\alpha_j)}$ 
  - $\alpha_i = 1.0$  for NUMA regions
  - $\alpha_i = \text{Perf}_{\text{GPU}}/\text{Perf}_{\text{CPU}}$  for GPU devices
- Within subtree:
  - NUMA partition: task-based factorization
  - GPU partition: batched factorization
- Root partition: use task-based factorization where tasks run either on a CPU core or in a GPU stream
- The DAG is dynamically scheduled in the root partition

# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level

# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level
- Two level parallelism in the DAG
  - Use OpenMP runtime system in [NUMA](#) regions as it comes with lower task management cost

# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level
- Two level parallelism in the DAG
  - Use OpenMP runtime system in [NUMA](#) regions as it comes with lower task management cost
- The dynamic scheduling is handled by the [Heteroprio](#)
  - Task scheduling criterion: acceleration factor on given resources

# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level
  - Two level parallelism in the DAG
    - Use OpenMP runtime system in [NUMA](#) regions as it comes with lower task management cost
  - The dynamic scheduling is handled by the [Heteroprio](#)
    - Task scheduling criterion: acceleration factor on given resources
- ▲ We are able to exploit GPU device while overcoming **memory limitations**



# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level
  - Two level parallelism in the DAG
    - Use OpenMP runtime system in [NUMA](#) regions as it comes with lower task management cost
  - The dynamic scheduling is handled by the [Heteroprio](#)
    - Task scheduling criterion: acceleration factor on given resources
- ▲ We are able to exploit GPU device while overcoming **memory limitations**
- ▲ Exploitation of **data locality** within each subtree

# Heterogeneous CPU-GPU multifrontal factorization

SpLDLT implementation details:

- We use [StarPU](#) to manage the DAG at the higher level
  - Two level parallelism in the DAG
    - Use OpenMP runtime system in [NUMA](#) regions as it comes with lower task management cost
  - The dynamic scheduling is handled by the [Heteroprio](#)
    - Task scheduling criterion: acceleration factor on given resources
- 
- ▲ We are able to exploit GPU device while overcoming **memory limitations**
  - ▲ Exploitation of **data locality** within each subtree
  - ▲ Reduce the impact of the cost for the CPU ↔ GPU memory transfers

## NUMERICAL EXPERIMENTS

---

# Numerical experiments

Experimental setup:

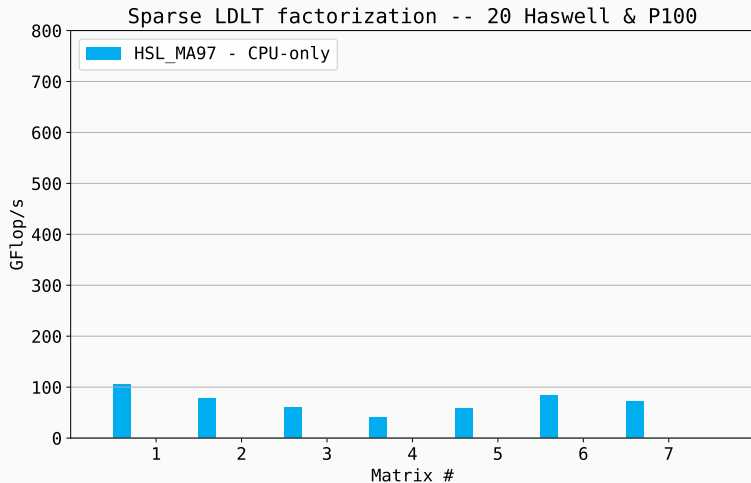
- 20 cores Intel Haswell machine (2 × E5-2695 v3)
- 1 NVIDIA Pascal P100
- 64 GB of RAM
- GNU compilers 7.1.0 with flags “-g -O2 -march=native”
- Intel MKL BLAS 11.3.1
- Metis 4.0.3

# Numerical experiments: Sparse $LDL^T$ factorization


#	Problem	$n$ $\times 10^3$	$nz(A)$ $\times 10^6$	$nz(L)$ $\times 10^6$	$flops$ $\times 10^9$
1	Oberwolfach/t3dh	79.17	2.22	50.60	70.10
2	Lin/Lin	256.00	1.01	126.00	285.00
3	PARSEC/H2O	67.02	2.22	234.00	1290.00
4	GHS_indef/sparsine	50.00	0.80	207.00	1390.00
5	PARSEC/Ge99H100	112.98	4.28	669.00	7070.00
6	PARSEC/Ga10As10H30	113.08	3.11	690.00	7280.00
7	PARSEC/Ga19As19H42	133.12	4.51	823.00	9100.00

- [HSL\\_MA97](#): multifrontal solver from the HSL library (CPU-only)
- [PARDISO](#): supernodal solvers part of the MKL library (CPU-only)
- [SpLDLT](#): our new solver, part of the SyLVER package (CPU-GPU)


# Numerical experiments: Sparse $LDL^T$ factorization



# Numerical experiments: Sparse $LDL^T$ factorization




# Numerical experiments: Sparse $LDL^T$ factorization



- CPU-only SpLDLT code is up to  $\times 3$  faster compared to PARDISO



# Numerical experiments: Sparse $LDL^T$ factorization




- CPU-only SpLDT code is up to  $\times 3$  faster compared to PARDISO
- Heterogeneous CPU-GPU SpLDT code is up to  $\times 3$  faster compared to CPU-only version

# Numerical experiments: Sparse Cholesky factorization


#	Problem	$n$ $\times 10^3$	$nz(A)$ $\times 10^6$	$nz(L)$ $\times 10^6$	$flops$ $\times 10^9$
1	Koutsovasilis/F1	344	13.6	173.7	218.8
2	Oberwolfach/boneS10	915	28.2	278.0	281.6
3	ND/nd12k	36.0	7.1	116.5	505.0
4	ND/nd24k	72.0	14.4	321.6	2054.4
5	Janna/Flan_1565	1565	59.5	1477.9	3859.8
6	Oberwolfach/bone010	987	36.3	1076.4	3876.2
7	GHS_psdef/audikw_1	944	39.3	1242.3	5804.1
8	Janna/Fault_639	639	14.6	1144.7	8283.9
9	Janna/Hook_1498	1498	31.2	1532.9	8891.3
10	Janna/Emilia_923	923	21.0	1729.9	13661.1

- **SpLDLT**: our new solver, part of the SyLVERpackage (CPU-GPU)

# Numerical experiments: Sparse Cholesky factorization



# Numerical experiments: Sparse Cholesky factorization



# Conclusions and future work

- Our sparse  $LDL^T$  factorization with APTP on multicore compares favourably with the state of the art solvers
- Using StarPU allows us GPU devices along with the multicores
- SpLDLT offers better numerical robustness compared to SBK like strategies
- We adapted the APTP pivoting to LU factorization..
- ..integrating it within SpLU using the same tree partitioning strategy as for SpLDLT

Thanks for listening!

Questions?